Lecture 7

Kernel Machines - Structured Prediction

~~BACE~~

BRACE → BRACE

Input

Output: must be a word!

→ only allow certain structured outputs

Gene sequences: GAGTCA → how protein structures

"the cat sat on the mat" → only allow grammatically correct sentences

Kernel Input can be structured ⇒ previous lecture

e.g., weighted degree kernels

Kernel Based Structured Prediction:

$$k : (\mathcal{X} \times \mathcal{Y}) \times (\mathcal{X} \times \mathcal{Y}) \to \mathbb{R}$$

PSD + symm.: Feature map $\phi : \mathcal{X} \times \mathcal{Y} \to \mathbb{R}^h$

Model: $f(y|x) = w^T \phi(x, y)$

Multi-class classification

$k((x,y),(x',y')) = \langle \phi(x,y), \phi(x',y') \rangle$ ← General

$K((x,y),(x',y')) = k(x,x') \mathbb{1}(y = y')$ ← for multiclass

$\mathcal{X} \in \mathbb{R}^d$, $\mathcal{Y} = \{1, \dots, C\}$

input (e.g. image) output (e.g. "dog")

Feature Map: $\psi(x,y) = \begin{pmatrix} \phi(x) \cdot \mathbb{1}(y=1) \\ \vdots \\ \phi(x) \cdot \mathbb{1}(y=C) \end{pmatrix}$

$$y|x = \underset{y}{\operatorname{argmax}} \langle w, \psi(x,y) \rangle = \underset{c}{\operatorname{argmax}} w_c^T \phi(x)$$

How to learn such a structured output model?

Model: $f(y|x) = w^T \phi(x,y)$ Prediction $y|x = \arg\max_{y \in Y} f(y|x)$

Find target-margin model by solving

$$\min_w \frac{1}{2}\|w\|^2 + C \sum_{n=1}^{N} \xi_n$$

Constraints: $\forall_{n=1}^{N}, \forall_{y \in Y_n}: \underbrace{w^T \phi(x_n, y_n)}_{A} - \underbrace{w^T \phi(x_n, y)}_{B} \geq 1 - \xi_n$

y_n: ground truth y: all possible evaluation

A: Score of ground truth B: score of all possible evaluation

A-B: difference of scores, $A - B \geq 1 \; (-\xi_n)$

ξ_n: Slack term: account for wrong labels, noise, ...

$\min_w \frac{1}{2}\|w\|^2$ → maximize margin

larger C → smaller ξ_n → hard margin constraint

Structured Prediction vs. HMM

HMM Pro: - unsupervised

- once model learned, procedure for prediction determined

Struct. Pred. Pro:

- feature map and loss function give more flexibility for structure of problem
- model parameter $w \in \mathbb{R}^h$ can be directly optimized for best performance on supervised task

Summary

- structured output learning enables prediction of structured outputs (trees, sequences)
- assign matching scores to input-output pairs
- supervised - kernel-based framework
- difficulty: efficiently infer which output $y \in Y$ maximizes the score $f(y|x)$